

Failure and integrity panel: limits are evidence, not exceptions

The evaluation reports method failures, neutral neural results, calibration gaps, bounded dataset scope, and closest-prior uncertainty instead of turning partial signals into final claims.

A. Robust aggregation can fail under label skew

Krum collapse

F1 0.0766

20% malicious attacks

Normal FPR

1.0000

Both attack types

Boundary: robust aggregation alone does not solve non-IID poisoning.



Evidence: `kaggle\_matrix\_v2\_summary.md`

B. Learned NumPy MLP gate is not supported

traffic_only

F1 0.0485


v0.3 MLP FedAvg

source_gate

F1 0.0209

Below traffic-only

Boundary: v0.3 is failure analysis, not a learned-gate success.



Evidence: `kaggle\_cti\_v03\_mlp\_summary.md`

C. Edge neural CTI gain is neutral/negative

traffic_only

F1 0.8241


Torch MLP v0.4

source_gate

-0.0065

vs traffic-only

Boundary: do not claim CTI/source evidence improves Edge neural IDS.



Evidence: `kaggle\_cti\_v04\_torch\_mlp\_summary.md`

D. Calibration remains unresolved through v0.8

traffic_only

ECE 0.0506

Scalar temperature ECE

cal_source_gate

ECE 0.1084

Group3 ECE: 0.1067

Boundary: F1/FPR is partial positive; calibration is not solved.



Evidence: `cti\_v08\_posthoc\_calibration\_summary.md`

E. Cross-domain evidence stays bounded

CICIoT2023

sampled / rare-family runs only

NetFlow

stress + per-family 100k cap

Do not write full raw-dataset final performance

Evidence: CICIoT, UNSW, NetFlow summaries and Table 3

F. Closest-prior claims are blocked

P17 FedKD-IDS

Information Fusion closest prior, metadata-only

Blocked claims

No absence claim; no superiority claim

Formal abstract gate remains closed

Evidence: claim registry; P17 comparison

Integrity rule for the manuscript

A result can strengthen the paper even when it is negative: failures define the contribution boundary, prevent overclaiming, and guide the next experiment instead of being hidden.